Kaustubh Hemant Sonawane

Bristol, UK | ksstubh.sonawane@gmail.com | +44 7388229743 | +91 8591441019 | www.github.com/Arishtanemi4 | www.linkedin.com/in/kaustubhson4 | https://kaustubh-sonawane-portfolio.vercel.app

UK Visa Status: Student Visa. Eligible for Graduate Visa (2 Years Full Working Rights)

16 October 2026

Hiring Team
Primer

Application: Senior Engineer, ML

Dear Primer hiring team,

I am applying for the Senior Engineer, ML role. Being one of the first ML engineers at Primer, building the function's foundations alongside the Staff Engineer and working out with product, data and engineering where ML can move the needle, is the shape of work I want to do next. For three years at HDFC Asset Management I owned machine learning and GenAI systems end to end in a regulated financial services setting: researching and testing the approach, designing the architecture, building the software, deploying it to real users and fixing it in production. I would like to bring that ownership to Primer.

At HDFC I designed and built the applications end to end: Python FastAPI microservices and a ReactJS frontend, retrieval with LlamaIndex and Azure OpenAI, a separate pgvector database per department, and agentic tool-calling where agents call Python functions to run SQL-backed comparisons and generate PDF reports on request. I packaged the services as Docker containers and shipped them on AWS (EC2 including a GPU instance, S3, ECR and Lambda) through a one-click Jenkins and GitLab CI/CD pipeline. Where the model choice was an engineering trade-off rather than a default, I self-hosted Llama-3 on a GPU EC2 instance for our largest department to control cost and kept the stronger Azure OpenAI outputs for the decision makers. I also led up to four interns as a de facto team lead, so the mentoring side of the role is familiar.

The closest work I have done to routing decisions is my MSc dissertation with AstraZeneca. I designed the desirability scoring and ranking algorithm that orders around 2,100 cancer cell lines by target inclusion and exclusion genes across six evidence layers, with a veto for exclusion genes and a calibration step that keeps scores comparable across lineages, and we refined the ranking with feedback from AstraZeneca scientists. Recommending a cell line from partial evidence is the same shape of decision as routing a payment, and I was careful to report what the ranking could and could not support rather than overstating it.

I want to be direct about where my experience differs from your list. My production work is Python services and GenAI rather than models trained and served with scikit-learn, XGBoost, TensorFlow or PyTorch; my deep learning is academic, the MSc multi-architecture study and a CNN project in my B.Tech. My experimentation has been taking systems from prototype to real users and changing them on what usage showed, not designed A/B tests with power analysis and confidence intervals, and I have not built multi-armed or contextual bandits - I understand where they apply and I would begin from a simple baseline before reaching for them. My monitoring is production troubleshooting and user feedback rather than an observability stack, my stakeholders have been internal business teams rather than Primer's merchants, and my experience is about three years rather than a long senior track record, though those years were spent owning systems end to end.

Primer has been remote-first since day one, which suits how I work; I am based in Bristol. I hold a Student Visa and am eligible for the Graduate Visa, which gives two years of full working rights. Thank you for considering my application; I would be glad to talk it through.

Yours sincerely,
Kaustubh Sonawane